

Lecture 1: Course Overview and DNN Basics

CS 580B1: Trustworthy Machine Learning

Fall 2024

Acknowledgments

Course slides derived from material created by Rajeev Alur and Osbert Bastani at UPenn ([link](#) to their course)

Other relevant courses:

[CS 329T](#) at Stanford

[CS 521](#) at UIUC

[ECE1784H/CSC2559H](#) at U Toronto


Agenda

- Course Overview
 - Motivation
 - Course Information
- DNN Basics
 - Computational Model
 - Training a DNN: Backpropagation

What is *trustworthy* machine learning?


Standard machine learning pipeline

Training / Learning a model



Goal: Maximize performance (e.g., accuracy) on training data


Testing / Evaluating a model



Goal: Evaluate performance (e.g., accuracy) on unseen data

Should we deploy a classifier that is 99% accurate on test data?


Lack of robustness



Lack of robustness



(A) **Cow: 0.99**, Pasture: 0.99, Grass: 0.99, No Person: 0.98, Mammal: 0.98




(B) No Person: 0.99, Water: 0.98, Beach: 0.97, Outdoors: 0.97, Seashore: 0.97




(C) No Person: 0.97, **Mammal: 0.96**, Water: 0.94, Beach: 0.94, Two: 0.94

Fig. 1. Recognition algorithms generalize poorly to new environments. Cows in ‘common’ contexts (e.g. Alpine pastures) are detected and classified correctly (A), while cows in uncommon contexts (beach, waves and boat) are not detected (B) or classified poorly (C). Top five labels and confidence produced by ClarifAI.com shown.

Lack of interpretability




(a) Husky classified as wolf



(b) Explanation

Figure 11: Raw data and explanation of a bad model's prediction in the "Husky vs Wolf" task.

Lack of interpretability



Zech et al., Variable generalization performance of a deep learning model to detect pneumonia in chest radiographs: A cross-sectional study, PLoS Med

Lack of alignment




Figure 4: Screenshots of harmful content generation from the examples shown in Figure 1: ChatGPT (top left), Claude 2 (top right), Bard (bottom left), LLaMA-2 (bottom right). Complete generations are shown in Appendix B.




Figure 1: Example of jailbreak prompt. Texts are adopted from our experimental results.

Lack of fairness


OPINION

MAY 18, 2023 | 5 MIN READ

Police Facial Recognition Technology Can't Tell Black People Apart

AI-powered facial recognition will lead to increased racial profiling

BY THADDEUS L. JOHNSON & NATASHA N. JOHNSON



Steffi Loos/Getty Images

Science

Current Issue First release papers Archive A

HOME > SCIENCE > VOL. 366, NO. 6464 > DISSECTING RACIAL BIAS IN AN ALGORITHM USED TO MANAGE THE HEALTH OF POPULATIONS

RESEARCH ARTICLE



Dissecting racial bias in an algorithm used to manage the health of populations

ZIAD OBERMEYER¹, BRIAN POWERS, CHRISTINE VOGELI, AND SENDHIL MULLAINATHAN² [Authors Info & Affiliations](#)

SCIENCE • 25 Oct 2019 • Vol 366, Issue 6464 • pp. 447-453 • DOI: 10.1126/science.aax2342

129,494 1,266



Racial bias in health algorithms

The U.S. health care system uses commercial algorithms to guide health decisions. Obermeyer *et al.* find evidence of racial bias in one widely used algorithm, such that Black patients assigned the same level of risk by the algorithm are sicker than White patients (see the Perspective by Benjamin). The authors estimated that this racial bias reduces the number of Black patients identified for extra care by more than half. Bias occurs because the algorithm uses health costs as a proxy for health needs. Less money is spent on Black patients who have the same level of need, and the algorithm thus falsely concludes that Black patients are healthier than equally sick White patients. Reformulating the algorithm so that it no longer uses costs as a proxy for needs eliminates the racial bias in predicting who needs extra care.

Science, this issue p. 447; see also p. 421

Lack of privacy



Figure 1: An image recovered using a new model inversion attack (left) and a training set image of the victim (right). The attacker is given only the person's name and access to a facial recognition system that returns a class confidence score.




Fig. 1: Membership inference attack in the black-box setting. The attacker queries the target model with a data record and obtains the model's prediction on that record. The prediction is a vector of probabilities, one per class, that the record belongs to a certain class. This prediction vector, along with the label of the target record, is passed to the attack model, which infers whether the record was *in* or *out* of the target model's training dataset.

Scalable Extraction of Training Data from (Production) Language Models

Milad Nasr^{*1} Nicholas Carlini^{*1} Jonathan Hayase^{1,2} Matthew Jagielski¹

A. Feder Cooper³ Daphne Ippolito^{1,4} Christopher A. Choquette-Choo¹

Eric Wallace⁵ Florian Tramèr⁶ Katherine Lee^{+1,3}

¹Google DeepMind ²University of Washington ³Cornell ⁴CMU ⁵UC Berkeley ⁶ETH Zurich

^{*}Equal contribution ⁺Senior author

Fredrikson et al., Model Inversion Attacks that Exploit Confidence Information and Basic Countermeasures, CCS'15

Shokri et al., Membership Inference Attacks against Machine Learning Models, IEEE SP'17

Lack of calibration




Figure 1. Confidence histograms (top) and reliability diagrams (bottom) for a 5-layer LeNet (left) and a 110-layer ResNet (right) on CIFAR-100. Refer to the text below for detailed illustration.

Measuring accuracy is not enough!

Beyond accuracy

Consider...

- A doctor using a model to diagnose a patient
- A judge using a model to decide a bail application
- The police department using a facial recognition model to recognize people to arrest
- A self-driving car using an object detection model detect obstacles and pedestrians

Does the model...

- Get fooled by small perturbations? (robustness)
- Make the decision for the right reasons? (interpretability)
- Discriminate against population subgroups? (fairness)
- Leak sensitive training data? (privacy)
- Correctly report its confidence? (calibration)

What is *trustworthy* machine learning?

- Techniques for evaluating *trustworthiness* of ML models and learning *trustworthy* models
- Metrics of *trustworthiness*:
 - Accuracy
 - Robustness
 - Interpretability
 - Alignment with human values
 - Fairness
 - Privacy
 - Being well-calibrated
 - ...
- Specific focus in this course on evaluation approaches that use *formal methods*

Course logistics: Communication

Course website - <https://www.cs.colostate.edu/~cs580b1/>

Includes syllabus, schedule, and course policies

[Canvas](#) for announcements, review submissions, and course materials including slides

[Teams](#) for discussion

Course logistics: Communication

Office hours – Wednesday, 4 PM – 5 PM and by appointment

Email at ravi.mangal@colostate.edu and include CS580B1 in subject

Anonymous feedback [here](#)

Course logistics: Grading break-up

Paper reviews	50%
Project	25%
Presentation	20%
Participation	5%

Course logistics: Paper reviews

- To be submitted by Wednesday, 11:59 PM
- Roughly 1-2 pages long
- Expected format for reviews will be shared
- Ideally, typeset using LaTeX

Course logistics: Presentation

- Presenters (team of 2 or 3) prepare a 30-minute presentation
- All students, including presenters, come ready with discussion points (which should be included in the paper reviews)

Course logistics: Project

- Proposal due by Oct 10
- Project options:
 - Literature review
 - Replicating a paper
 - Applying studied techniques to your own ML models
 - Exploring new research idea
 - (talk to me)
- Final presentations in the week before finals

Course logistics: Workload

- 1 review every week
- 2 paper presentations / lead discussions over the semester
- Project proposal, final report, and presentation

DNN Basics

- A simple feedforward neural network with 1 hidden layer:

$$f_{\theta}(x) := W^{(2)} \left(g(W^{(1)}x + b^{(1)}) \right) + b^{(2)}$$

where

- $f_{\theta}: \mathbb{R}^i \rightarrow \mathbb{R}^o$,
- $W^{(1)} \in \mathbb{R}^{k \times i}$, $b^{(1)} \in \mathbb{R}^k$, $W^{(2)} \in \mathbb{R}^{o \times k}$, $b^{(2)} \in \mathbb{R}^o$ are parameters (or weights) to be learnt and k is the **number of hidden neurons**,
- $x \in \mathbb{R}^i$ is an input,
- $g: \mathbb{R} \rightarrow \mathbb{R}$ is an **activation function** applied component-wise, i.e.,
$$g \left(\begin{bmatrix} z_1 \\ z_2 \end{bmatrix} \right) = \begin{bmatrix} g(z_1) \\ g(z_2) \end{bmatrix}$$